

Exploiting `Math.expm1(-0)` in v8 TurboFan JIT Compiler

Ryan Torok, Sara McFearin, Michael Jarrett, Rachel Shi

December 7, 2019

1 Introduction

Browser bugs are difficult to find, but they appear to be prevalent across all four major browser engines. In recent years most of the focus has been on bugs in the JavaScript Just-in-Time (JIT) compilers. Improper optimization can often lead to a memory corruption exploit which allows the attacker to take control of the victim's browser and begin executing arbitrary code on the victim's machine, making browser bugs a popular target for cybercriminals who run botnets, ransomware scams, and more. Although it is not uncommon for compilers to contain bugs in their large codebases, browser JIT compilers are unique in the sense that they must deal with adversarially chosen code. In a normal pre-compilation setting, if a compiler bug is found, programmers simply won't write code that triggers the bug for sake of security. However, in browsers, the compiler runs on the user's machine, and if a browser JIT compiler bug is found, attackers will intentionally ship code which triggers the bug in order to take over the user's machine.

With this in mind, it seems like we will be playing an infinite game of Whack-a-mole with browser JavaScript engines, but new techniques have risen in recent years to make bug elimination faster, most notably Fuzzing. Fuzzing is a technique which originated in image encoding protocols, where a penetration tester will pass random binary inputs to the protocol attempting to cause a crash. With some modifications based on knowledge of how browser JITs create a graph of a code segment, penetration testers can write automated tools which generate random JavaScript code inputs which look somewhat interesting. The most popular of these tools is called *Fuzzily*, which has already found a plethora of bugs in the JavaScript engines for all four major web browsers.

2 The Bug

We will focus now on one of these bugs in particular. The TurboFan JIT compiler used by v8, the JavaScript engine for the Chrome and Chromium web browsers, was exploited in 2018 using an edge case with the function `Math.expm1()`, which computes $e^x - 1$ for argument x . Specifically, if we evaluate `Math.expm1(-0)`, this should produce the value

```

418 Type OperationType::NumberExpml(Type type) {
419     DCHECK(type.Is(Type::Number()));
420     return Type::Union(Type::PlainNumber(), Type::NaN(), zone());
421 }
422

```

Figure 1: The buggy return type declaration for `Math.expml()` in `operation-typer.cc`

```

let isNegativeZero = Object.is(Math.expml(-0), -0);|

```

Figure 2: Code which determines whether `Math.expml()` returns 0 or -0

-0. However, the TurboFan JIT lists the return types for this function as a union of the `PlainNumber` type and `NaN`. This union includes all values of a 64-bit floating point number, except -0. V8 defines this behavior using a special table in `typer.cc` and `operation-typer.cc`. The code from the latter is shown in Figure 1. The JIT uses this fine-grained type information to perform variable range analysis that is used in array bounds check eliminations. For example, if TurboFan realizes a boolean variable is used to index an array of length ≥ 2 , then the native compiled code can forgo ensuring the array index is in range, saving time. As we will see in the following section, the mismatch between the expected and actual output range of `Math.expml()` can have catastrophic effects.

3 Exploitation Techniques

In this section we will walk through the process of exploiting the bug, all the way up to arbitrary code execution. In our instance, we choose to spawn a shell.

3.1 Triggering the Bug

The goal of the first stage of our exploit is to utilize the buggy fine-grained return type of `Math.expml()` to eliminate a bounds check which is in fact not safe, and allow us to access memory outside the bounds of the array. The first step involves creating a variable which depends on the result of our buggy function call. In other words, we need a way to distinguish between the values 0 and -0. The only function useful to us in this case is `Object.is()`. In the code segment shown in Figure 2, the variable `isNegativeZero` will be 1 if `Math.expml()` returns -0, and 0 otherwise. Running this line of code before it becomes 'hot' and JIT optimized will always store 1 into `isNegativeZero`, because the code is interpreted directly and the native backing function for `Math.expml()` returns -0.

This becomes interesting when we run the code after it has been JIT-optimized. Consider the code in Figure 3. This code calls the function `hax` once without optimization, and once with optimization. The `%OptimizeFucntionOnNextCall` annotation is a macro which

```
function hax(x) {
    return Object.is(Math.expm1(x), -0);
}

console.log(hax(-0));
%OptimizeFunctionOnNextCall(hax);
console.log(hax(-0));
```

Figure 3: The first call to `hax()` correctly prints 'true', but the second prints 'false' because the buggy JIT optimization folds the entire `Object.is` call to 'false'.

```
function hax(x) {
    var victim = [0.1, 0.2, 0.3, 0.4]
    var a = Object.is(Math.expm1(x), -0);
    return victim[a * 1234];
}

console.log(hax(-0));
%OptimizeFunctionOnNextCall(hax);
console.log(hax(-0));
```

Figure 4: First attempt at an out of bounds array access. This fails because the `Object.is` call folds to false, which prevents the out of bounds index.

can be invoked using the `--allow-natives-syntax` flag on the command line. These annotations are only used for debugging, and in a real exploit this macro would be replaced with calling the function a set number of times in a loop to achieve the level of optimization we want. Running this code prints `true` for the first call and `false` for the second. Why? The first time the code runs, it is interpreted directly, and the engine correctly evaluates `-0` equal to `-0`. Before the second call, the JIT (incorrectly) optimizes the function to native code, and in its analysis it incorrectly assumes `Math.expm1()` can never return `-0`, since its return type is (erroneously) listed as a union of `PlainNumber` and `NaN`. Therefore TurboFan assumes the `Object.is` call always evaluates to `false`, and uses *Folding* to reduce the variable `isNegativeZero` to the constant `false`, thus breaking the semantics of the `Math.expm1()` function.

3.2 Triggering an Out-of-Bounds Array Access

Now that we've triggered the bug, how can we turn it into a memory leak? The idea is to exploit the bug to cause the JIT to remove an array bounds check in the optimized code, allowing us to read out of bounds. Consider the code in Figure 4

4 Patch and Resolution